

Canadians should use memory

Pierre Bergé^{*,1}, Julien Hemery², Arpad Rimmel², and Joanna Tomasik²

¹LRI, Université Paris-Sud, Université Paris-Saclay, 91405 Orsay Cedex, France

²LRI, CentraleSupélec, Université Paris-Saclay, 91405 Orsay Cedex, France

Abstract

We establish that the competitive ratio of any randomized memoryless strategy is not less than $2k + 1$. Only randomized strategies using memory potentially overpass this ratio now.

1 Definitions

We start by introducing the notation. For any graph $G = (V, E, \omega)$, let $G \setminus E'$ denotes its subgraph $(V, E \setminus E', \omega)$. If P is a path, we note its cost as $\omega(P) = \sum_{e \in P} \omega(e)$.

1.1 Memoryless Strategies for the k -CTP

We remind the definition of CTP. Let $G = (V, E, \omega)$ be an undirected graph with positive weights. The objective is to make a traveller traverse the graph from a source node s to a target one t , with $s, t \in V$. There is a set $E_* \subsetneq E$ of blocked edges. The traveller does not know a priori which edges are blocked. He discovers a blocked edge only when arriving to one of its endpoints. For example, if (v, w) is a blockage he will discover it when arriving to v (or w). The goal is to design the strategy A with the minimal competitive ratio.

We focus on *memoryless strategies* (MS). Concretely, we suppose that the traveller remembers the blocked edges he has discovered but forgets the nodes which he has already visited. In other words, a decision of an MS is independent of the nodes already visited. In the literature, the term *memoryless* was used in the context of online algorithms (e.g. PAGING PROBLEM [?], LIST UPDATE PROBLEM [?]) which take decisions according to the current state, ignoring past events. An MS can be either deterministic or randomized.

Definition 1 (Memoryless Strategies for the k -CTP) *A deterministic strategy A is an MS if and only if (iff) the next node w the traveller visits depends on graph G deprived of blocked edges already discovered E'_* and the current traveller position v : $w = A(G \setminus E'_*, v)$. Similarly, a randomized strategy A is an MS iff node w is the realization of a discrete random variable $X = A(G \setminus E'_*, v)$.*

MSes are easy to be implemented because they do not use past moves to take a decision. For the same reason, they do not need any extra memory either.

For example, the GREEDY strategy [?] is an MS. It consists in choosing at each step the first edge of the shortest path between the current node v and the target t . This strategy, illustrated

*Corresponding author: Pierre.Berge@lri.fr

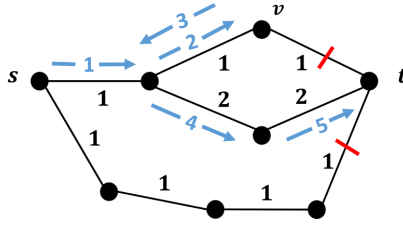


Figure 1: Illustration of the GREEDY strategy on a graph with 2 blocked edges

in Figure 1, does not refer to anterior moves. In contrast, the REPOSITION strategy [?] is not an MS as any decision refers to the past moves of the traveller.

We propose the following process to identify whether a strategy A is a deterministic MS. Let us suppose that a traveller T_1 executes strategy A : he has already visited certain nodes of the graph, he is currently at node v but he has not reached target t yet. Let us imagine a second traveller T_2 who is airdropped on node v and starts applying strategy A . If the traveller T_2 always follows the same path as T_1 until reaching t , A is a deterministic MS. If T_1 and T_2 may follow different paths, then A is not an MS. Formally, proving that a strategy is an MS consists in finding the function which transforms the pair $(G \setminus E_*, v)$ into node $w = A(G \setminus E'_*, v)$.

1.2 Competitive ratio

Let (G, E_*) be a *road map*, i.e. a pair with graph $G = (V, E, \omega)$ and blocked edges $E_* \subsetneq E$, such that there is an (s, t) -path in graph $G \setminus E_*$ (nodes s and t remain in the same connected component when all blocked edges are discovered). We note $\omega_A(G, E_*)$ the distance traversed by the traveller reaching t with strategy A on graph G with blocked edges E_* and $\omega_{\min}(G, E_*)$ the cost of the shortest (s, t) -path in graph $G \setminus E_*$.

The ratio $\omega_A(G, E_*) / \omega_{\min}(G, E_*)$ is abbreviated as $c_A(G, E_*)$. A strategy A is c_A -competitive [?, ?] iff for any (G, E_*) , $\omega_A(G, E_*) \leq c_A \omega_{\min}(G, E_*)$. Otherwise stated, for any (G, E_*) , $c_A(G, E_*) \leq c_A$. If strategy A is randomized, $\omega_A(G, E_*)$ is replaced by $\mathbb{E}(\omega_A(G, E_*))$ which is the expected distance traversed by the traveller to reach t with strategy A . The competitive ratio can also be evaluated on a family \mathcal{R} of road maps, put formally,

$$c_{A, \mathcal{R}} = \max_{(G, E_*) \in \mathcal{R}} c_A(G, E_*) \quad (1)$$

This “local” competitive ratio fulfils $c_{A, \mathcal{R}} \leq c_A$. The competitive ratio can also be extended to families of strategies. We note c_{MS} the competitive ratio of MSes, which is the minimum over competitive ratios of any MSes: $c_{\text{MS}} = \min_{A \text{ MS}} c_A$. In the remainder, we identify a set of road maps \mathcal{R} such that the competitive ratio of MSes on it is $2k + 1$, i.e. $c_{\text{MS}, \mathcal{R}} \geq 2k + 1$.

1.3 Graphs G_k

We define recursively a sequence of graphs G_i for $i \geq 1$ with weights from $\{1, \varepsilon\}$, $0 < \varepsilon \ll 1$. Graphs G_1 and G_{i+1} are represented in Figures 2a and 2b, respectively. Edges with weight 1 are thicker than edges with weight ε . For every graph G_k , we will only consider the road maps where the blocked edges are at the right of the graph, the left edges affecting negligibly the total cost of a traveler. We note E_{right} the set of all the right edges.

When the traveller discover blocked edges in G_k , parts of the graph become *dead ends*. If a traveller visits a dead end, the only chance for him to reach t is to leave it anyway. Dead ends




Figure 2: Recursive construction of graphs G_i

only provide to the traveller additional costs which make the competitive ratio increase. As a consequence, the most competitive strategies will not visit dead ends. From now on, as soon as a dead end appear when the traveller discovers a blockage, we make it disappear from the graph: graph $G \setminus E'$ becomes graph G deprived of both edges E' and dead ends appeared when edges E' are deleted. Let \mathcal{G}_k denote the set of all sub-graphs of G_k where at most k edges have been removed. Formally, $\mathcal{G}_k = \{G_k \setminus E' : |E'| \leq k\}$.

1.4 Equivalent binary trees

We define for each graph of \mathcal{G}_k an equivalent binary tree, which will represent all changes of direction from t . We note T_\emptyset the empty tree, and (v, T_a, T_b) a non empty tree with $v \in V$, T_a the binary tree above and T_b the binary tree below.

We obtain the equivalent binary tree by applying the BIN-TREE algorithm. This algorithm take as an entry a directed tree obtained by first removing the left part of the graph and then directed all edges from t . Then, the BIN-TREE algorithm will construct a binary tree according to the outdegree of each node.

Algorithm 1: BIN-TREE algorithm

Data: a directed tree T and a node v

Result: binary tree

if $\deg^-(v) = 0$ **then return** T_\emptyset ;

if $\deg^-(v) = 1$ **then return** BIN-TREE(T , v_{next}) ;

if $\deg^-(v) = 2$ **then return** (v , BIN-TREE(T , v_{up}), BIN-TREE(T , v_{down}) ;

Given an edge in a binary tree, we will note $A(e)$ the child edge from above and $B(e)$ the child edge from below. We also note $C(e)$ the set of all the descendants of e . Formally, $C(e) = C(A(e)) \cup C(B(e)) \cup \{A(e), B(e)\}$. Finally, we note $P(e)$ the parent edge of e : $P(e) = e' \Leftrightarrow e = A(e')$ or $e = B(e')$.

We define recursively a set of edges E_k , E_0 being the edges connected to the root and $E_k = \{e : P(e) \in E_{k-1}\}$. E_k represents the edges which are at distance k from the root.

When an edge is cut from the original graph, the equivalent tree can be easily obtain by replacing the tree (v, T_a, T_b) by T_a (resp. T_b) if the cut edge is before T_b (resp. T_a)

Lemma 1 For any road map $(G, E_*) \in \mathcal{R}_k$, $E_{T_G}(j-1) = 2^j$ with $j = |E_*|$

Proof. If $G = G_k$, then $|E_*| = k$ and the binary tree is complete, so $E_{T_{G_k}}(k-1) = 2^k$.

We suppose now for any road map with $|E_*| = j$ that $E_{T_G}(j-1) = 2^j$. Let (G, E_*) a road map with $|E_*| = j-1$ Then it exists a road map $(G', E_* \cup \{e\})$ with G obtained by removing e to G' .

First, we remark that for any binary tree $T = (v, T_a, T_b)$ and i , if the number of edge of depth i is 2^{i+1} , the tree is complete until depth i , so T_a and T_b are complete until depth $i - 1$ and each have 2^i edges at depth $i - 1$: $|E_T(i)| = 2^{i+1} \implies |E_{T_a}(i - 1)| = |E_{T_b}(i - 1)| = 2^i$.

Second, we remark that for any binary tree T and i , if the number of edge of depth i is 2^{i+1} , then the number of edge for the depth $i - 1$ is 2^i : $|E_T(i)| = 2^{i+1} \implies |E_T(i - 1)| = 2^i$.

Let $T' = (v, T_a, T_b)$ the sub-tree of T'_G where v is the root of the edge e (we will suppose e is on the T_a side). Then if e is at depth i then the remarks and the hypothesis give $|E_T(j - i - 1)| = 2^{j-i}$ and $|E_T(j - i - 2)| = 2^{j-i-1} = |E_{T_a}(j - i - 2)|$. By cutting e , T become in G the tree T_a . We find that the number of edge at the depth $j - i - 2$ remains. So the rest of the tree being unchanged by cutting e , the number of edges of T_G for depth $j - 2$ is unchanged from T'_G and is equal to 2^{j-1} : $|E_{T_G}(j - 2)| = 2^{j-1}$. ■